

100 AI UGC in under an hour

The full playbook, built for UGC

Hey, it's Tiv. You commented "UGC", here's the full playbook. Basically it breaks down exactly what we do at the agency to pump out our UGC, but **step by step, by hand**, like you had none of it wired into a skill. You'll see everything under the hood: the stack, the file setup, the brief, the prompts to write, the editing, the scaling. You'll understand exactly how the machine runs.

One thing to get right away: the "**100 UGC in an hour**" doesn't come from doing all this by hand, one video at a time. It comes from having this process wired into a skill that runs in parallel. So it's your call: grab ours to ship your 100 right now, or build your own with this guide. Real talk, 90% of people will never build their own, and that's fine.

100+

UGC per session

<1h

start to finish

0

editing software

~\$2

per UGC (real cost)

What you'll produce

By the end of a session: varied UGC scripts (talking head, storytelling, before/after, demo...), realistic people talking about your product (voice + lipsync), word-level captions, a punchy edit with music, A/B variants ready to test, and a production log to keep track. 9:16 for TikTok or 1:1 / 4:5 for Meta, ready to upload.

1. Prerequisites (install once)

These tools don't change between projects. Install once, they work for every future session. 10 minutes max.

1. Claude Code — it runs everything

Claude's command line. It writes the scripts, calls the generation models, drives FFmpeg, handles the files. Without it, nothing chains together.

```
# Requires Node.js 18+
npm install -g @anthropic-ai/claude-code
# then, inside your working folder:
cd ~/my-ugc && claude
```

2. FFmpeg — the video engine

Open-source, free, not a SaaS. It cuts, assembles, captions and exports. Claude writes the commands, FFmpeg runs them. 10x faster than software with a UI.

```
# macOS
brew install ffmpeg
# Ubuntu / Debian: sudo apt install ffmpeg · Windows: ffmpeg.org + PATH
```

3. A pay-per-use generation API key

To generate the videos of people talking. No subscription: you pay as you go, ~\$0.13 per generated second. One key covers image and video.

Install order: Node → Claude Code → FFmpeg, then your key in a **.env** file (never in plain text in the chat, never on Git).

2. Organize your files + the brief

Claude doesn't guess where your files are. With the right structure, it knows where everything is from second one and doesn't waste time.

```
my-ugc/
|-- assets/
|   |-- product/      # packshots, product photos
|   |-- footage/      # your own clips if you have them
|   '-- music/        # royalty-free .mp3
|-- scripts/          # scripts generated by Claude
|-- output/
|   |-- final/        # finished UGC
|   '-- variants/     # A/B variants
|-- .env              # your API key
'-- brief.txt         # THE most important doc
```

brief.txt is the doc Claude reads first. It's not a prompt, you fill it in plain language. The more precise you are, the better the output.

PRODUCT: [exact name + what it does in 1 line]
AVATAR: [who buys: age, style, pains, desires]
ANGLES: [2-3 marketing angles to run]
CASTING: [who talks: e.g. woman ~24, brown hair, casual]
SCENE: [bedroom / car / bathroom / living room / outside]
TOPE: [energetic, complicit, no bullshit...]
FORMAT / DURATION: [9:16 TikTok, 15-20s]
OFFER / CTA: [e.g. -30% this weekend / link in bio]
WHAT TO AVOID: [robotic voice, too "ad-like", white background]

The avatar is the most important field. "Everyone" = 0 hook that lands. The tighter you go, the more it converts.

3. The UGC pipeline, step by step

Here's the full flow, in order. **Each blue block is a prompt you'd have to write and tune yourself** if you do it all by hand. In the skill, these prompts are already written and up to date: we just talk to Claude and it does the work. Here, you see the raw grind.

Step 1 — Audit ~1 min

First, Claude takes inventory so it doesn't assume anything.

PROMPT 1 — AUDIT

Analyze the contents of my /assets/ folder and read /brief.txt. Summarize the brief in 5 points, list the available assets (product, footage, music), and tell me what's missing to produce quality UGC.

Result: you instantly know if a music track or a packshot is missing, and fix it before spending.

Step 2 — Generate the scripts ~3 min

Claude reads the brief and writes full UGC scripts (not ideas, spoken scripts, shot by shot).

PROMPT 2 — SCRIPTS

Based on the brief, generate 10 varied UGC scripts. For each: the TYPE (talking head, storytelling, before/after, demo, shock hook...), the target DURATION, the HOOK (first 3 seconds), the SCRIPT spoken out loud (contractions, natural tone, one personal tangent), and the CTA. Save to /scripts/script_01.txt through script_10.txt.

→ Hook first, because that's what makes or breaks a creative. Saved files = you tweak without re-prompting.

Type	Hook	Duration	Platform
Talking head	Authentic to-camera	15–30s	TikTok, Meta
Storytelling	A question or a problem	30–60s	Meta Feed, YouTube
Before / After	Problem visible instantly	15s	TikTok, Meta
Shock hook + reveal	Stat or surprising fact	7–10s	TikTok
Demo	Real-time use	15–30s	Meta, Google

Step 3 — Generate the talking heads ~15 min (parallel)

This is the heart of UGC. From a starting image (the keyframe) and the script, a reference-to-video model outputs a realistic person who **talks**: voice and lipsync in one pass. Two prompts: one for the image, one for the animation.

PROMPT 3A — THE KEYFRAME (STARTING IMAGE)

Generate an image: [casting from the brief] in [scene], shot selfie-style, raw iPhone look, natural light, realistic colors, skin with grain, no filter, no beauty retouching. Tight face framing. Generate 3 variants, I'll pick the most believable.

PROMPT 3B — THE ANIMATION (THE PERSON TALKS)

Animate this image: the person says exactly "[script]". Tone [energetic/complicit]. She [precise gesture, e.g. points at the product, tucks a strand of hair]. Camera nearly static. One hand always holds the camera. Natural delivery, not recited.

The golden rule: you dictate everything (what she says, the tone, the gesture). If you let the model improvise, it does whatever (she says "look at this" and shows nothing). Generate the exact duration of the script, no more (every extra second = money).

Step 4 — Word-level captions ~2 min

The audio gets transcribed to get the timing of each word. Claude generates a styled caption file and burns it in with FFmpeg.

```
# word-level transcription -> styled .ass file -> burn-in
ffmpeg -i clip.mp4 -vf "ass=captions.ass" -c:a copy sub.mp4
```

Result: captions land right on the words, native TikTok style, never off.

Step 5 — The edit (cuts, silences, music) ~5 min

Claude cuts the silences for a tight rhythm, adds punch-ins, lays the music under the voice, exports at your platform's format.

```
# 1. detect silences > 0.5s
ffmpeg -i sub.mp4 -af silencedetect=noise=-30dB:d=0.5 -f null -
# 2. cut the gaps + music under the voice + 9:16 export
ffmpeg -i sub.mp4 -i music.mp3 -filter_complex \
"[1:a]volume=0.15[m];[0:a][m]amix=inputs=2[a]" \
-map 0:v -map "[a]" -c:v libx264 -b:v 8M -r 30 -s 1080x1920 final.mp4
```

Result: a finished video, punchy, right format, ready to post.

Step 6 — A/B variants ~10 min

For each finished UGC, Claude auto-generates variants (different hook, different music, different format) and generates a CSV so you find your way around the Ads Manager.

PROMPT 6 — VARIANTS

For each UGC in /output/final/, generate: 2 alternate hooks (question version + direct benefit version), 1 music variant, and the 1:1 format for Meta Feed. Name [name]_hA / _hB / _mB / _1l. Generate a CSV: file, variant type, difference.

4. Level up: b-rolls + your own footage

Two ways to go further than a plain talking head.

Generate product b-rolls to cut in

To break up a talking head and make it livelier, cut in generated product shots at the moments the script names the product.

PROMPT — PRODUCT B-ROLLS

Generate 3 product b-rolls from /assets/product/: a close-up on the detail, a slow rotation on a neutral background, a lifestyle setup. 5s each, natural light, no text. Cut them in where the script names the product.

If you already have your own footage

No need to generate everything. If you've filmed (you, a creator, product clips), Claude edits for you, and there it's \$0 in API.

PROMPT — EDITING YOUR FOOTAGE

Here's my footage in /assets/footage/. Write a voiceover from the brief, lay it over the shots, cut the dead time, add word-level captions and music under the voice, export in 9:16. Give me 3 hook variants.

Cost: if you just edit your footage without generating anything, it's \$0 in API. You only pay when Claude generates video.

5.15 UGC hooks (ready to adapt)

The hook is 80% of your creative's fate. Pick one, swap the [brackets] for your product, test 4 per script.

1. "I tested [product] for 30 days, here's what happened"
2. "Nobody told you this about [problem]"
3. "Stop [common mistake] right now"
4. "The thing [niche] doesn't want you to know about [topic]"
5. "I wish I knew this before buying [category]"
6. "POV: you find [product] and you can't go back"
7. "3 signs you need [product]"
8. "This [product] at \$[price] beats [alternative] at 3x the price"
9. "If you have [problem], watch this till the end"
10. "The [product] I've used every day for [duration]"
11. "Before / after [duration] with [product]"
12. "You do [mistake]? Here's why it's not working for you"
13. "I spent \$[X] on [category] to tell you which one's worth it"
14. "Nobody talks about what [product] actually does"
15. "Stop scrolling if you have [problem]"

6. Realism: what separates real from slop

The tech, Claude handles. The real challenge is credibility. Most AI UGC gets spotted in a second. The principles that make it read real:

- **Raw iPhone, never "pretty".** Phone look, natural colors, zero filter. A too-polished video = spotted instantly.
- **Intent in every gesture.** She looks at the right spot, shows what she names. Dictated, never improvised.
- **Some imperfection.** A natural micro-gesture, a hesitation, a personal tangent. That's what humanizes it.
- **The right pace.** Not recited, not robotic. She talks like she's talking to a friend.

Good UGC	Bad UGC
✓ Hook targeting a precise pain	✗ Generic hook ("discover this product")
✓ Raw look, shot on a phone	✗ Too clean, plastic render
✓ Word-level captions, concise	✗ A whole paragraph on screen
✓ A clear CTA	✗ No CTA

7. Scale: from 10 scripts to 100 UGC

10 well-built scripts give 100+ testable creatives. You decline, you don't create new. Each UGC multiplies across these dimensions:

Dimension	Options	×
Hook (first 3s)	Question / Shock / Direct benefit / Before-after	×4
Music	Hype / Chill / Emotional / TikTok trending	×4
Format	9:16 / 1:1 / 4:5 / 16:9	×4
Casting	Run the same mechanic on 3-5 different people	×5
Language	EN / FR / DE / ES / IT	×5

Production runs in parallel. You tell Claude "run all generations in parallel and the FFmpeg exports with 4 workers max." That's what makes 100 UGC fit in an hour.

8. The test strategy (in what order)

Phase	Nb	What you test	Budget/day
1. Hook test	8-12	1 script, 4 hooks, 2-3 music. Find the hook that grabs.	\$5-10
2. Concept test	6-8	3-4 scripts with the winning hook. Find the concept that converts.	\$10-20
3. Scale	3-5	The 2-3 best, all variants. Maximize ROAS.	\$50+

The simultaneous-test trap. Testing 50 creatives at once gives no clear signal. Start with 8-12, find the winning hook (CTR > 2%), then scale on that one only.

9. Managing costs

Scenario	Nb UGC	Est. cost	Detail
Quick test	10	~\$15-25	Short talking heads, no b-roll
Standard batch	50	~\$80-120	15-20s UGC + variants
Full session	100	~\$150-250	All variants, several castings

2 optimizations: generate short previews first to validate the concept before rendering the full duration, and reuse validated keyframes instead of regenerating. A failed test caught early = money saved.

10. Troubleshooting (the 8 common errors)

Render looks "AI" / too clean. Cause: beauty-cam look. Fix: force the raw iPhone in the keyframe (neutral colors, grain, zero retouching) and generate 3 candidates, keep the most believable.

She says "look" but shows nothing. Cause: gesture left to the model. Fix: dictate the precise gesture in the animation prompt, timed to the word.

Captions slightly off. Cause: constant-speed timing. Fix: lock the captions to the real transcription (word by word), not proportionally.

Different character from one clip to the next. Cause: no fixed reference. Fix: keep the same keyframe (the same starting image) for all clips of one UGC.

Voice too "studio" / metallic. Cause: the native voice comes out too clean. Fix: run the audio through a "phone" filter (light highpass + lowpass + a small room reverb) before export.

FFmpeg: "Invalid data" or truncated text. Fix: re-encode to standard MP4 H.264 (ffmpeg -i in.mp4 -c:v libx264 -c:a aac out.mp4) and use the absolute font path.

Audio out of sync between clips. Cause: different audio sample rates. Fix: "resample all audio to 44100 Hz before assembling."

Generation too slow. Cause: Claude generates sequentially. Fix: "run all generations in parallel and the exports with 4 workers max."

The shortcut

You just saw everything you'd have to write, tune and fix to make UGC by hand: the prompts, FFmpeg, the casting, the misses to clean up.

Honestly, you can do it yourself. You can even take the time to wire it all into your own skill and adapt the prompts so they work broad on your brand. But I won't lie to you: it's way more complex than it looks, and it's really not obvious to set up. Count several weeks and a good bit of testing before it actually runs.

Us, we've already done it, and that's the difference: **no prompt to write**. You talk to Claude like a friend, it already knows this whole guide by heart. You install it in 2 minutes, you say what you want, you get your UGC. We update it every month and run it every day on dozens of brands.

[Access the Creative System →](#)

Meanwhile, take this playbook and get started. The best time to start is now.